

Operation Analytics and Investigating Metric Spike



Introduction:

Operational Analytics plays a pivotal role in enhancing organizational efficiency by dissecting the entirety of a company's operations. Businesses can use this analytical method as a compass to identify areas that are ready for optimization and improvement. Your work as a data analyst is varied and requires you to work closely with various departments, including marketing, operations, and support. Utilizing the abundance of data that these teams have gathered and turning it into useful insights is your main goal.

Project Description:

Understanding and improving a company's end-to-end operations depend heavily on operational analytics. Working with a variety of departments, such as operations, support, and marketing, is the main goal of a data analyst's job to extract useful insights from data. This helps the business make well-informed decisions and optimize its processes.

Examining abrupt spikes in metrics is a crucial component of operational analytics, and it calls for a thorough grasp of data analysis methods. For this project, you will be working for a fictitious corporation that is similar to Microsoft as a Lead Data Analyst. You have to use your sophisticated SQL skills to examine the datasets that have been provided, providing operational improvement advice and insight into metric changes.

Approach:

A. Jobs Reviewed Over Time:

- Objective: Calculate the number of jobs reviewed per hour for each day in November 2020.
- Your Task: Write an SQL query to calculate the number of jobs reviewed per hour for each day in November 2020.

The screenshot shows a SQL IDE with a query editor and a result grid. The query is as follows:

```
1 -- Write an SQL query to calculate the number of jobs reviewed per hour for each day in November 2020
2 • use job_db;
3
4 • select * from job_data;
5
6 • select ds, count(job_id) as count_jobs, sum(job_id)/3600 as reviewed_in_hours from job_data
7   where ds>='2020-11-01' and ds<='2020-11-30'
8   group by ds
9
10
11
12
13
14
```

The result grid shows the following data:

ds	count_jobs	reviewed_in_hours
2020-11-30	2	0.0119
2020-11-29	1	0.0064
2020-11-28	2	0.0133
2020-11-27	1	0.0031
2020-11-26	1	0.0064
2020-11-25	1	0.0056

B. Throughput Analysis:

- Objective: Calculate the 7-day rolling average of throughput (number of events per second).
- Your Task: Write an SQL query to calculate the 7-day rolling average of throughput. Additionally, explain whether you prefer using the daily metric or the 7-day rolling average for throughput, and why.
The daily throughput

job_data creation

Jobs Reviewed Over Time

Throughput Analysis*

Language Share

Limit to 1000 rows

4

5 • `select ds,count(event)/sum(time_spent)`

6 `from job_data;`

7

8

9 • `select ds,count(event)/sum(time_spent) as throughput`

10 `from job_data`

11 `group by ds`

12

13

14

Result Grid		Filter Rows:		Export:		Wrap Cell Content:	
ds	throughput						
2020-11-30	0.0500						
2020-11-29	0.0500						
2020-11-28	0.0606						
2020-11-27	0.0096						
2020-11-26	0.0179						
2020-11-25	0.0222						

And the highest being 0.0606

Rolling metrics serve as a powerful tool for monitoring and interpreting daily fluctuations in key performance indicators. By smoothing out short-term variations and highlighting longer-term trends, they offer valuable insights into the trajectory of your metrics, empowering you to make informed decisions and drive continuous improvement.

C. Language Share Analysis:

- Objective: Calculate the percentage share of each language in the last 30 days.
- Your Task: Write an SQL query to calculate the percentage share of each language over the last 30 days.

job_data creation Jobs Reviewed Over Time Throughput Analysis **Language Share Analysis*** Duplicate Rows Detection*

Limit to 1000 rows

```

1 • use job_db;
2
3 • select * from job_data;
4
5 • select language, count(language) as count_lang, (count(language)/total_lang)*100 as percent_lang
6   from job_data
7   inner join(
8     select count(language) as total_lang from job_data) as inner_table
9   on 1=1
10  where ds>='2020-11-01' and ds<='2020-11-30'
11  group by language, inner_table.total_lang
12  order by percent_lang desc;
13
14

```

Result Grid Filter Rows: Export: Wrap Cell Content: [IA](#)

	language	count_lang	percent_lang
▶	Persian	3	37.5000
	English	1	12.5000
	Arabic	1	12.5000
	Hindi	1	12.5000
	French	1	12.5000
	Italian	1	12.5000

Here is the percentage of language used. With the highest being 37.50 by 'The Persian' language

D. Duplicate Rows Detection:

- Objective: Identify duplicate rows in the data.
- Your Task: Write an SQL query to display duplicate rows from the job_data table

job_data creation Jobs Reviewed Over Time Throughput Analysis Language Share Analysis* **Duplicate Rows Detection***

Limit to 1000 rows

```

14
15 -- duplicate rows in actor_id
16
17 • select * from job_data
18   where actor_id in(
19     select actor_id from job_data
20     group by actor_id
21     having count(*)>1);
22
23
24
25
26
27

```

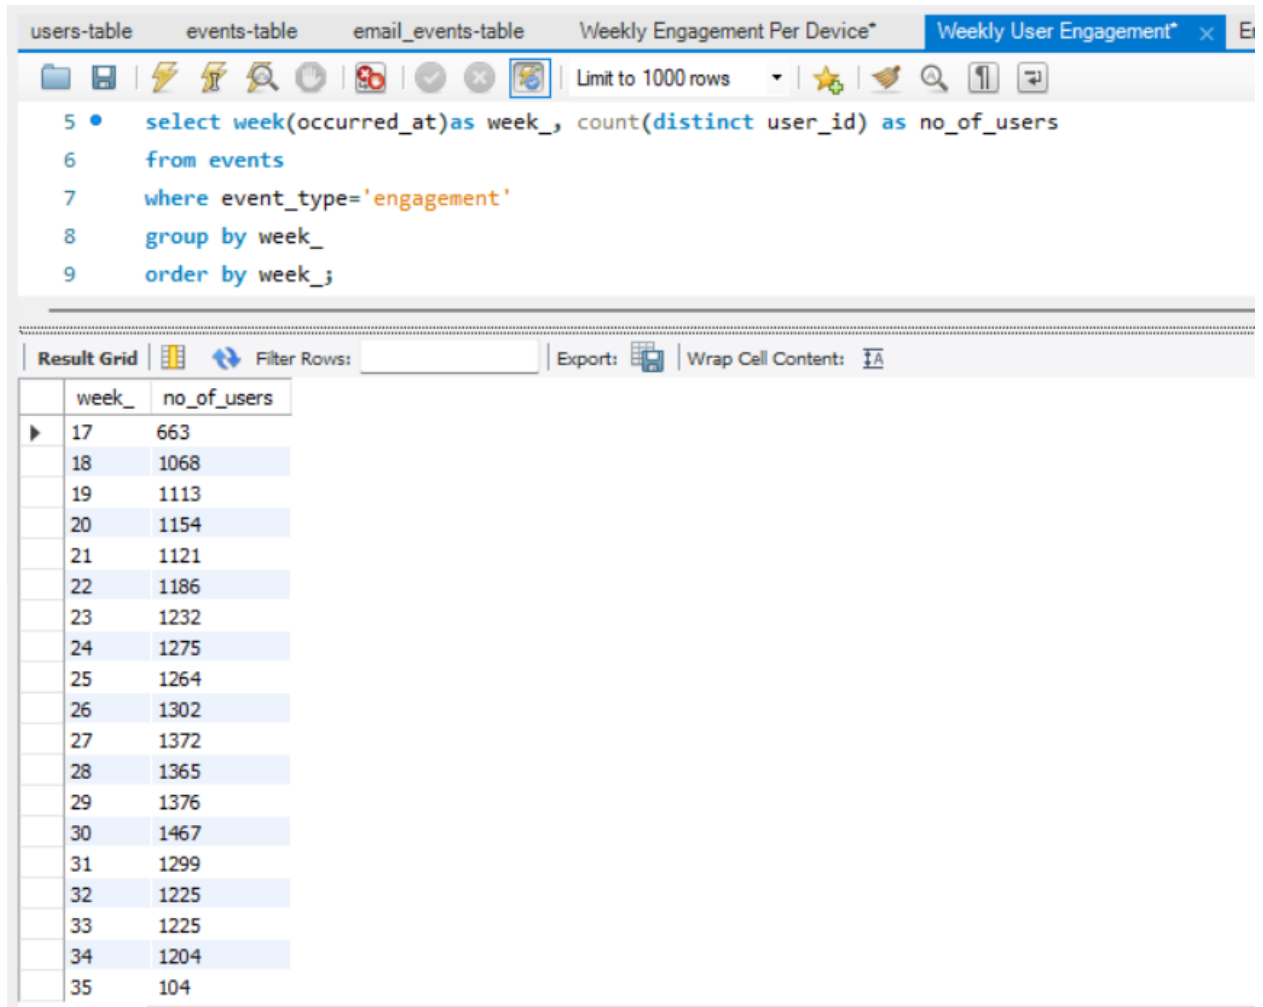
Result Grid Filter Rows: Export: Wrap Cell Content: [IA](#)

	ds	job_id	actor_id	event	language	time_spent	org
▶	2020-11-29	23	1003	decision	Persian	20	C
	2020-11-25	20	1003	transfer	Italian	45	C

Actor id 1003 has 2 duplicate rows and here are the full details of those duplicate rows

Case Study 2: Investigating Metric Spike

A. Weekly User Engagement:



users-table events-table email_events-table Weekly Engagement Per Device* Weekly User Engagement* x E

Limit to 1000 rows

```
5 • select week(occurred_at) as week_, count(distinct user_id) as no_of_users
6   from events
7   where event_type='engagement'
8   group by week_
9   order by week_;
```

Result Grid Filter Rows: Export: Wrap Cell Content:

	week_	no_of_users
▶	17	663
	18	1068
	19	1113
	20	1154
	21	1121
	22	1186
	23	1232
	24	1275
	25	1264
	26	1302
	27	1372
	28	1365
	29	1376
	30	1467
	31	1299
	32	1225
	33	1225
	34	1204
	35	104

Highest No_of_users in a week: 30

Lowest No_of_users in a week: 35

B. User Growth Analysis:

users-table | events-table | email_events-table | Weekly Engagement Per Device* | Weekly User Engagement* | Email Engagement A

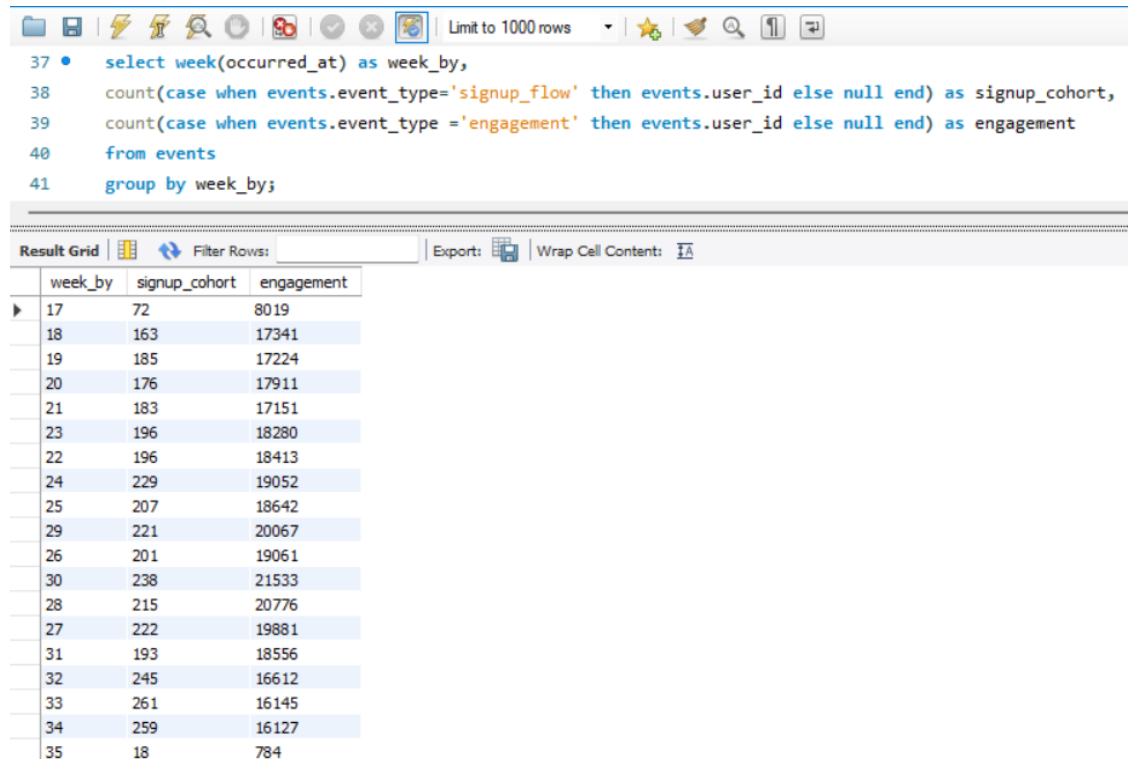
The lowest was on week 35 of the year 2014 with no_of_users as 18 and the highest was on week 33 of the year 2014 with no_of_users 261.

C. Weekly Retention Analysis:

Objective: Analyze the retention of users on a weekly basis after signing up for a product.

Your Task: Write an SQL query to calculate the weekly retention

Here's the week by retention



The screenshot shows a SQL IDE interface. At the top, there's a toolbar with various icons and a dropdown menu set to 'Limit to 1000 rows'. Below the toolbar, a SQL query is entered in a text area. The query is as follows:

```
37 • select week(occurred_at) as week_by,  
38        count(case when events.event_type='signup_flow' then events.user_id else null end) as signup_cohort,  
39        count(case when events.event_type='engagement' then events.user_id else null end) as engagement  
40    from events  
41    group by week_by;
```

Below the query editor, there's a 'Result Grid' section. It includes a 'Filter Rows:' input field, an 'Export:' button, and a 'Wrap Cell Content:' checkbox. The result grid displays a table with three columns: 'week_by', 'signup_cohort', and 'engagement'. The table contains 19 rows of data, with the first row highlighted in blue. The data is as follows:

	week_by	signup_cohort	engagement
▶	17	72	8019
	18	163	17341
	19	185	17224
	20	176	17911
	21	183	17151
	23	196	18280
	22	196	18413
	24	229	19052
	25	207	18642
	29	221	20067
	26	201	19061
	30	238	21533
	28	215	20776
	27	222	19881
	31	193	18556
	32	245	16612
	33	261	16145
	34	259	16127
	35	18	784

D. Weekly Engagement Per Device

```

7 • select concat(year(occurred_at),'-',week(occurred_at)) as week_by,device,count(distinct user_id) as weekly_eng
8 from events
9 where event_type='engagement'
10 group by week_by,device
11 order by week_by;

```

week_by	device	weekly_eng
2014-17	acer aspire desktop	9
2014-17	acer aspire notebook	20
2014-17	amazon fire phone	4
2014-17	asus chromebook	21
2014-17	dell inspiron desktop	18
2014-17	dell inspiron notebook	46
2014-17	hp pavilion desktop	14
2014-17	htc one	16
2014-17	ipad air	27
2014-17	ipad mini	19
2014-17	iphone 4s	21
2014-17	iphone 5	65
2014-17	iphone 5s	42
2014-17	kindle fire	6
2014-17	lenovo thinkpad	86
2014-17	mac mini	6
2014-17	macbook air	54
2014-17	macbook pro	143
2014-17	nexus 10	16

Week 30 of the year 2014 has had the highest user engagement of 322 users, the device being used was “MacBook Pro” for the week.

E. Email Engagement Analysis

users-table | events-table | email_events-table | Weekly Engagement Per Device* | Weekly User Engagement* | Email Engagement Analysis*

Limit to 1000 rows

```
23
24
25
26
27
28 • select (count_open/email_sent)*100 as per_open,(count_click/email_sent)*100 as per_click
29 from(
30   select count(case when action='email_open' then 1 end) as count_open
31   from email_events) as open_table
32 join(
33   select count(case when action='email_clickthrough' then 1 end) as count_click,
34   count(case when action in('sent_weekly_digest', 'sent_reengagement_email') then 1 end) as email_sent
35   from email_events) click_table;
36
37
```

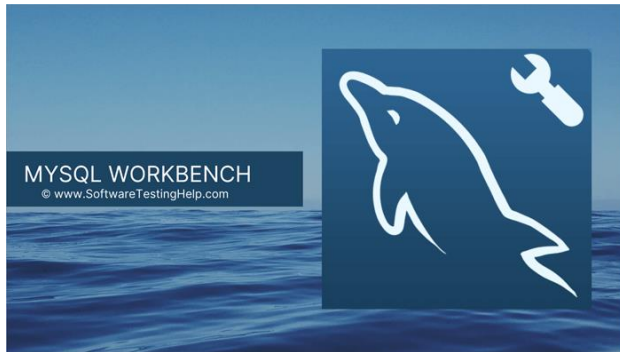
Result Grid | | Filter Rows: | Export: | Wrap Cell Contents:

per_open	per_click
33.5834	14.7899

Out of the total emails sent 33.5834% of the emails were opened and 14.7899% were clicked
 This query calculates the percentage of emails opened and clicks based on the total number of emails sent.

Tech-Stack Used:

For this project, MySQL Workbench is utilized for data analysis. MySQL Workbench provides a user-friendly interface for querying databases, making it suitable for executing complex SQL queries. Its functionalities allow for efficient data manipulation, extraction, and visualization, facilitating a comprehensive analysis of the datasets provided.



Insights:

Through the analysis, several key insights and observations can be drawn from the data. These include patterns in user behavior, trends in metric fluctuations, and correlations between different operational variables. Insights gained from the analysis can shed light on areas for improvement, identify potential bottlenecks, and highlight opportunities for optimization within the company's operations.

Result:

This project helped me with SQL Proficiency and also made me more comfortable with writing complex SQL queries to extract valuable insights from diverse datasets. And I have also learned a lot about Time-Series Analysis, Investigating metrics over periods, such as hourly, daily, and weekly. The project aims to achieve a comprehensive understanding of the company's operations through data-driven insights. By successfully executing the provided tasks and generating meaningful analyses, the project contributes to informed decision-making processes within the organization. Ultimately, the insights gleaned from the data analysis endeavor to enhance operational efficiency and drive business growth.

Conclusion:

This operational analytics project has been a worthwhile and rewarding experience that has provided a thorough immersion into the complexities of assessing a company's end-to-end operations. I have explored a variety of datasets, addressed a wide range of analytical issues, and produced useful insights to promote business progress throughout the project. I have improved our ability to query, aggregate, and analyze data by using advanced SQL techniques.

We have also become more adept at data visualization, time-series analysis, and problem-solving.